

Multi-Agent Clinical Intelligence System


Doctor Dashboard v3.1 • PDF + Voice Clinical Intelligence System

Agent Architecture Data Flow Skills Framework Safety & Security

Agent Architecture Flow

Orchestrator → Agents → Skills → Tools → LLM

Current runtime: the Express server supports **two intake paths**. PDFs go through **DocumentTypeRouter**, which classifies the document and routes it to a specialized extractor agent. Voice dictations go through a **swappable STT boundary**, then **VoiceExtractorAgent**, shared dashboard mapping, and **dashboard-readiness validation** before a record can remain processed. In the current workspace, Gemini STT is the active path; the voice architecture is intentionally modular so **Whisper can become the primary STT backend** with Gemini as fallback. **ReactExtractionAgent** exists in the repository but is optional rather than the default document-processing path. Processing Insights are served separately from analytics.sqlite through /api/analitics/overview.




Data Flow Architecture

Four operational flows, one storage boundary, and a single grounded path back to the clinician




Skills Framework

PDF, voice, presentation, validation, and chat capabilities organized by category



Tools Layer (PDF, voice, chat, and presentation utilities)



Tools are low-level utilities used by skills. Skills compose tools into reusable capabilities.

Safety & Security Architecture

Multi-layer protection for clinical AI systems

